

AI APPRENTICE

PROLOG From the Bottom Up

Apprentice: One who is learning by practical experience under skilled workers a trade, art or calling.

—Webster's Seventh New Collegiate Dictionary

"In the beginner's mind there are many possibilities, in the expert's there are few."

—S. Suzuki, Zen Mind, Beginner's Mind

Since AI is such a new field, very few have had the opportunity of the apprentice to study under a master craftsman. Except for those working in large university and corporate AI labs, most of us who are interested in AI are working in relative isolation, hungrily devouring all of the books, articles, and inexpensive software we can get our hands on.

In this column we would like to become your colleagues in the exchange of ideas, techniques, and information about resources. We hope that as we share our ideas with you, you will do the same with us. If you have a better solution to a problem than the one we present, let us know. Also, if you have a problem that is of general interest, or an unusual technique, by all means send it along.

We are planning to use actual code in presenting new programming techniques. In the column itself we will print chunks of code, and at times we will make programs available on the bulletin boards listed on page 7. Since we're gearing this column to those who, like us, come from the more traditional programming backgrounds, we've decided to work primarily in Pascal. This seems like a good choice not only because it's a good teaching language but also because it's fa-



Our focus is
on how PROLOG
operates and what
distinguishes it from
conventional
programming
languages

miliar to most C and Modula programmers. As with everything else, our direction will be heavily influenced by your opinions, so be sure to be free with them.

We've always had the idea that one of the most effective ways to learn a new programming language is to write a program that implements part of the language. In our first two columns, we will describe a

very basic implementation of PROLOG. Our focus in the first article is not how to program in PROLOG but how PROLOG operates and what features distinguish it from more conventional programming languages. In the next column we will discuss the programming structures that are used to implement these features. This will not only provide you with a valuable look inside the language but also demonstrate some interesting Pascal programming techniques.

You adventuresome souls who want to get a jump on the next column may want to download the completed VT-PROLOG (very tiny) interpreter written in Turbo Pascal, which is available on the bulletin boards.

A SIMPLE EXAMPLE

We will begin to try to understand how PROLOG operates by examining a simple example. Consider a collection of sentences. The sentences consist of single lowercase letters, called facts, or more complicated structures, called rules. The rules consist of a single lowercase letter and a special symbol (:-), followed by a fact or series of facts separated by commas. Each sentence is terminated by a period. This collection of sentences is called a data base. The following is a sample data base:

1. a :- b,c,d.
2. a :- e,f.
3. c.
4. d.
5. e.
6. f.
7. g :- e,f.

The numbers aren't part of the

data base, they are merely for our convenience. Line 1 can be read "a is true if b is true and c is true and d is true." Line 3 says "c is true." The *as*, *bs*, etc., can stand for objects in the real world, but since PROLOG programs manipulate symbols and not their meaning, we will not confuse the issue by attaching an explicit meaning to each symbol. In line 1, *a* is called the head and *b,c,d* is called the tail of the rule. Line 3 has a head but no tail. The commas in rules stand for the word "and."

This data base may not seem particularly useful, but it is still possible to perform some interesting operations upon it. For instance, we might like to ask some questions about its contents. We can see that *c* is true within the data base, but what about *a* and *e*? Are facts *a* and *e* both true within this data base? A question about the data base is called a query. We write this query as *a,e*. It is easy to examine the data base and see if *a,e* is true, but as the data base becomes larger and the queries become more complicated, a systematic method of searching the data base will be necessary. We

The indentation level represents the level of recursion

can develop this method by noticing a couple of facts. First, an empty query (symbolized by *NIL*) is always true. Second, if the head of the query and the head of the rule match, replacing the head of the query by the tail of the rule doesn't change whether the query is true or false.

The second statement may require a little explanation. To prove the query *a,e*, we notice that line 1 of the data base says that *a* is true if *b,c*, and *d* are true. In other words, *a,e* is true if *b,c*, and *d* are true and *e* is true. That condition can be expressed by the query *b,c,d,e*. If we can find a series of transformations such as this that reduce the query to *NIL*, then the original query was true.

One problem still remains with this method. Some transformations may lead to dead ends. For example, *a,e* = > *b,c,d,e* will fail because no statement in the data base be-

LISTING 1. A procedure to solve queries.

```
PROCEDURE solve(query) ;
  VAR
    i : integer ;
  BEGIN
    IF query = NIL
      THEN write('yes')
    ELSE
      FOR i := 1 TO max_rule_number DO
        IF head(rule[i]) = head(query)
          THEN solve( append(tail(rule[i]),tail(query)) ) ;
  END ; (* solve *)
```

gins with *b*, so no further transformation is possible. In a case like this we back up to the query before the transformation and try the next rule. We only say that the query was false when all possible transformations of the original query have been exhausted. The process of backing up and trying a new rule is called backtracking.

We can express all this a little more concisely with the pseudocode procedure shown in Listing 1. This pseudocode looks a great deal like Pascal, but we don't want to worry about things like data types for rules and queries yet so we will ignore them for the present.

In Listing 1, *rule[i]* is simply one line from the data base. *Head* is a function that returns the first item of either a sentence or a query. *Tail* returns everything in a sentence or query after the head. If there is nothing following the head, *tail* returns *NIL*. *Append* is a procedure that merges pieces of rules and queries to produce a new query. For example, *head(rule[1]) = a*, *tail(rule[1]) = b,c,d*, and *append(tail(rule[1]),tail(query)) = b,c,d,e*. Appending *NIL* to a query returns the original query: *append(f,NIL) = append(NIL,f) = f*. The *solve* procedure

is recursive; it calls itself with the transformed query. The recursion is terminated when either a query has been reduced to *NIL* or the search of the data base for a particular query is exhausted.

To get a feeling for this process, let's examine the solution to the query *a,e* in some detail. Listing 2 shows the calls to *solve* and the transformation of the queries. The indentation level represents the level of recursion.

UNIFICATION

One reason for collecting facts in a data base is to represent relationships among the data items and to use those relationships to answer questions about the data. To accomplish this, we will first introduce a more convenient notation to represent rules and queries and then modify the *solve* procedure to handle the new notation.

Let's consider a relatively simple data base, which contains some information about a few people's personal preferences:

1. likes(joan,pool).
2. likes(alice,candy).
3. female(joan).
4. female(alice).

LISTING 2. Calls to procedure *solve* to answer query ?- *a,e*.

call to procedure	matching rule
solve(a,e)	1. a:- b,c,d
solve(b,c,d,e)	no matching rule
solve(a,e)	2. a:- e,f
solve(e,f,e)	5. e
solve(f,e)	6. f
solve(e)	5. e
solve(NIL)	write 'yes', recursion terminated
solve(e)	no matching rule (backing out of recursion we continue search)
solve(f,e)	no match
solve(e,f,e)	no match
solve(a,e)	no match, finished

5. male(paul).
 6. likes(paul,X) :- likes(X,pool), female(X).

Again, the numbers aren't part of the data base. Readers familiar with PROLOG may recognize this notation. It is a convenient way to represent relationships. Line 1 may be read "Joan likes pool." "Likes" is called a functor and names a relationship between the first component, "Joan," and the second, "pool." Line 3 says "Joan is a female."

Terms like "Joan", "female," etc., are loaded with millions of associations in our minds, but here we will only be manipulating symbols. Thus the knowledge that Joan is a common female name in North America has no meaning unless it is defined in the data base.

Line 6 is slightly different from the others: it contains a variable. In PROLOG's notation, constants begin with lowercase letters and variables begin with uppercase letters or an underscore. A variable is a symbol whose value will be determined in the context of some query. Line 6 can be read "Paul likes X if X likes pool and X is female." Or,

The process of finding a variable's value is called binding the variable

more simply, "Paul likes any female who likes pool."

With this notation we form queries such as "Does Paul like Joan?" or "Who does Paul like?" The former is expressed as `?- likes(paul,joan)` while the latter is expressed as `?- likes(paul,A)`. `?-` indicates that what follows is to be treated as a query. The first query will cause `solve` to respond "yes" or "no." In the second case, `solve` should not only tell us that Paul likes someone but who that someone is. The process of finding a value for a variable is called binding the variable, and the value of the variable is often called its binding. The set of all the current variables and their bindings is called the environment.

Of course, with the introduction of functors and variables, our previous version of `solve` is inadequate.

The problem lies in the process of matching the head of a rule against the head of a query. Consider the query `?- likes(paul,joan)`. The head of the query is `likes(paul,joan)`. We take the functor and its components paul and joan to be a single entity. We can probably agree that it should not match the head of rules 1 through 5. Rule 1 tells us about what joan likes, so the functors match but the first components don't. Similarly, rule 2 tells us what alice likes so it doesn't match the query. Rules 3, 4, and 5 aren't about likes at all so they don't match the query. Rule 6 is different, however. The functors match and the first components match, but what about joan and X? In a case such as this, we say that the head of rule 6 and the query match, with the variable X from the rule bound to joan.

Now we can continue as before. We append the tail of the query `NIL` with the tail of the rule `likes(X,pool),female(X)`, and try to solve this transformed query along with the added information that X is bound to joan. The new query could be interpreted as `likes(joan,pool),female(joan)`.

COMING SOON ...

An Expert System that will offer more power and flexibility than artificial intelligence programs costing ten times more.

An Expert System that will open up new mass markets for artificial intelligence.

An Expert System at an unbelievable Paperback price!

EXCLUSIVE AAAI SNEAK PREVIEW

Paperback Software will be demonstrating this powerful and unique tool at the American Association of Artificial Intelligence Conference in Philadelphia. August 11-15, 1986, Booth #509.

DON'T MISS IT!!

**From
PAPERBACK SOFTWARE**

Paperback Software is a registered trademark of Paperback Software International. • 2830 Ninth Street, Berkeley, CA 94710



Suppose instead we had posed the query `?- likes(paul,A)`. Rules 1 through 5 still don't provide us with a match but the head of rule 6 will match, provided we can bind `X` to `A`. When one variable is bound to another, we say that the variables share a binding. When one is matched to a constant, the other automatically takes on the same value.

We still have one complication to consider. What would happen if we were to add a new rule to the data base, for example, `likes(joan,X) :- likes(X,candy)`? Are the `X`s in this rule the same as in rule 6? Or suppose we had posed the query as `?- likes(paul,X)`. Is the `X` in the query the same one that appears in the rules?

Questions like this are really an informal way of asking what the scope of a variable is. A variable's scope is the range over which its bindings are valid. Readers who are familiar with block-structured languages like Pascal or C understand the concept of scope. In a block-structured language a variable's scope is the block in which it is defined. This makes it possible to define a local variable, like `X`, in a pro-

LISTING 3. Solve procedure using unification.

```
PROCEDURE solve(query,env,level) ;
  VAR
    i : integer ;
    new_env : same as query and env ;
  BEGIN
    IF query = NIL
      THEN print_env(env)
      ELSE
        FOR i := 1 TO max_rule_no DO
          IF unify(copy(head(rule[i])),level+1),head(query),env,new_env)
            THEN solve(append(copy(tail(rule[i])),level+1),tail(query),
                        new_env,level+1) ;
    END ; (* solve *)
```

Questions like this are really an informal way of asking what the scope of a variable is

cedure and also to define another local variable with the same name in another procedure. Since these variables have different scopes, they are different variables.

In PROLOG, a variable's scope is the rule or query that contains it.

This means that in `likes(paul,X) :- likes(X,pool),female(X)`, each instance of `X` represents the same variable. That variable is entirely different from the `X` defined in `?- likes(paul,X)`. In order to avoid confusion among variables with the same name but entirely different meanings, we will have `solve` make a copy of the rule being examined. The copy will be exactly the same as the original rule in the data base but all variables in the rule will be tagged by appending the recursion level to them. Since a rule either fails or causes a recursive call to `solve`, this will mean

NGS FORTH

*79 STANDARD

FORTH 79 with double number standard extensions serves as a base for the NGS FORTH system. All 79-Standard words are available for truly transportable applications.

*NGS ENHANCEMENTS

NGS FORTH provides a variety of words found to be useful by FORTH programmers (including a print spooler).

FIG LOOKALIKE MODE

Portability from your previous FIG FORTH applications is assured through a well defined FIG compatibility mode.

*MS-DOS COMPATIBLE

NGS FORTH operates on top of MS-DOS providing complete access to all DOS interrupts and functions. Regular DOS text files may be written, read, or loaded with ease. Hardware I/O port reading and writing are also supported.

*ON-LINE CONFIGURABLE

NGS FORTH features on-line reconfiguration for expansion or contraction of all memory areas. This includes dictionary, data stack, return stack, user variable, and buffer memory locations and sizes.

*ENVIRONMENT SAVE AND LOAD

An ultra-fast environment store and reload saves time when using large vocabularies.

*MULTI-SEGMENTED

NGS FORTH is multi-segmented to support large applications. Separate code, data, and stack segments can each grow to a 64K byte maximum.

*EXTENDED ADDRESSING

An extended addressing pointer can access data anywhere within the full megabyte 8088 addressing range. A high speed inter-segment move word transfers data rapidly anywhere in memory.

*AUTO LOAD SCREEN BOOT

NGS FORTH can be configured to auto-load any disk screen at startup.

*LINE AND SCREEN EDITORS

The full screen editor utilizes the special keys of the PC keyboard. The standard (FIG) line editor is also included.

*DECOMPILER AND DEBUGGING AIDS

The NGS decompiler provides immediate decompilation of FORTH colon definitions. The locator reports the screen number from which a FORTH word was loaded. The tracer provides real time monitoring of an application's execution.

*8088 ASSEMBLER

A basic 8088 assembler is provided with control structures and standard mnemonics.

*BASIC GRAPHICS AND SOUND

Graphics words are defined for many of the more useful operations: point setting, line drawing, area filling, and color. A wide range of sounds can be generated with the sound words.

*VERY FAST

NGS FORTH is a very fast FORTH optimized for the IBM PERSONAL COMPUTER and MS-DOS compatibles.

***PRICES START AT \$70**

***NEW HP-110 & HP-150 VERSIONS AVAILABLE**



NEXT GENERATION SYSTEMS
P.O. BOX 2987
Santa Clara, CAL. 95055

CIRCLE 27 ON READER SERVICE CARD

ITEM IN THE RULE	ITEM WITHIN THE QUERY	CONSTANT (C2)	VARIABLE (V2)	FUNCTOR (F2)
Constant (C1)	succeed if C1 = C2	succeed bind V2 to C1	succeed	fail
Variable (V1)	succeed bind V1 to C2	succeed bind V1 to V2	succeed bind V1 to F2	succeed if expressions have same functor and arity and each pair of components can be unified
Functor (F1)	fail	succeed bind V2 to F1		

TABLE 1. Unification of items in rules and queries.

that when seeking the resolution of a query, each time a rule is encountered its variables will have unique names. Listing 3 contains the new version of *solve*.

Solve is called with the current query as before. Two additional parameters are included with each call: *env*, the current environment, and *level*, an integer that contains the current recursion level. *Env* is a list of variables and their bindings. *Print_env* is a routine that prints the current environment. If the query is resolved (if some set of transformations reduce it to *NIL*), the current environment can be printed to show what set of bindings led to the resolution of the query. *Copy* is a function that is passed an object to copy and an integer. It then returns a copy of that object with the integer appended to each of the variables in the object.

Unify is a Boolean function that matches the heads of sentences against the head of the query. If it returns a value of true (indicating that a match has been made), *new_env* contains a copy of the old environment plus any new bindings that *unify* may have made. (The routine is named *unify* rather than something like *match* because the process of matching and binding is formally called unification. *Unify* isn't particularly difficult to design. We'll leave aside most of its details, but Table 1 indicates how rules and queries should be unified.

Unify will have to be recursive be-

cause of the requirement to unify each of the components of a complex expression. Arity, mentioned in the table, is the number of components in the expression.

This description of the interpreter is somewhat concise, so let's look at a few steps in the resolution of the query *?- likes(paul,A)*. To get a better understanding of this process, Table 2 shows the steps involved in resolving this query. One thing to note in this table is that

Unify is a Boolean function that matches the heads of sentences against the head of the query

when *unify* encounters a variable it looks up that variable's binding in the current environment. If it finds a binding, it attempts to unify the binding with the parameter being matched. Thus, in attempting to unify *joan* with *X#0*, it will look up the binding of *X#0* and then try to unify *joan* and *A*. Since *A* isn't bound, these two items are unified by binding *A* to *joan*.

Although this description has been greatly simplified and does not consider such important concepts as how arithmetic and negation are handled, we hope it gives you an understanding of how sim-

ple relationships can be placed in a data base and then retrieved. If you download the program, experiment with setting up a simple relationship data base such as the one we've discussed. This will help you become more accustomed to looking at heads and tails and understanding the process of instantiation that makes PROLOG such a powerful symbol manipulator.

In the next issue we'll jump right in and see how to implement VT-PROLOG. AI

Bill and Bev Thompson are writers and consultants specializing in implementing AI techniques on microcomputers. They are the authors of MicroExpert, an expert system shell, and have worked extensively on knowledge-based designs.

Resources

The literature of PROLOG is growing rapidly. The basic form of the VT-PROLOG interpreter comes from an article "Describing PROLOG by its Interpretation and Compilation" by Jacques Cohen, December 1985, *Communications of the ACM*. In the same issue is the article "PROLOG in 10 Figures" by Alain Colmerauer, one of the original developers of PROLOG. This article describes PROLOG by means of a series of figures. The figures are both eye-catching and informative. "Learning about PROLOG," a useful introductory article on PROLOG programming by Ramachandran Bharath and Margaret Sklar, appeared in the July 1985 issue of *COMPUTER LANGUAGE*, pp. 49-54. The August 1985 issue of *BYTE* was dedicated to declarative languages, of which PROLOG is the most important example.

PD PROLOG is a public domain PROLOG interpreter available on a variety of bulletin board systems and was released by Automata Design Associates, 1370 Arran Way, Dresher, Pa. 19025.

TABLE 2. Solving the query *?- likes(paul,A)*.

QUERY	ENVIRONMENT	LEVEL	MATCHING RULE	BIND
likes(paul,A)	NIL	0	6 likes(paul,X) :- likes(X,pool), female(X)	X#0 to A
likes(X#0,pool),female(X#0)	(X#0 A)	1	1 likes(joan,pool)	A to joan
female(X#0)	(A joan) (X#0 A)	2	3 female(joan)	
NIL				